

A. Prediction Performance

Classifier	Clean A / F1 / Conf.	Grayscale A / Δ / C	Hue A / Δ / C	Translation 32 px A / Δ / C	Patch 4×4 A / Δ / C
ResNet head	0.976 / 0.976 / 0.936	0.960 / -1.6 pp / 0.966	0.934 / -4.2 pp / 0.938	0.966 / -1.0 pp / 0.987	0.898 / -7.8 pp / 0.916
ViT head	0.972 / 0.972 / 0.962	0.956 / -1.6 pp / 0.954	0.942 / -3.0 pp / 0.944	0.963 / -0.9 pp / 0.975	0.910 / -6.2 pp / 0.922
CLIP head	0.984 / 0.984 / 0.614	0.952 / -3.2 pp / 0.962	0.932 / -5.2 pp / 0.938	0.975 / -0.9 pp / 0.981	0.798 / -18.6 pp / 0.798
CLIP zero-shot	0.968 / 0.968 / 0.943	0.944 / -2.4 pp / 0.954	0.926 / -4.2 pp / 0.932	0.956 / -1.2 pp / 0.973	0.780 / -18.8 pp / 0.784

B. Cue-Conflict Decisions

Classifier	Shape / Texture / Other	Content Accuracy / F1	Coverage	Shape bias
ResNet head	93 / 22 / 85	0.465 / 0.448	0.575	0.809
ViT head	138 / 14 / 48	0.690 / 0.696	0.760	0.908
CLIP head	145 / 3 / 52	0.725 / 0.711	0.740	0.980
CLIP zero-shot	145 / 5 / 50	0.725 / 0.718	0.750	0.967

C. Representation Cosine Similarity

Backbone	Grayscale	Hue	Translation 32 px	Patch 4×4	Cue conflict Content / Style
ResNet-50	0.789	0.759	0.933	0.712	0.358 / 0.202
ViT-B/16	0.739	0.717	0.942	0.650	0.359 / 0.061
CLIP ViT-B/32	0.895	0.892	0.969	0.735	0.727 / 0.611